need.

- **True negatives**—Model predicts negative, actually negative. This is good.
- **True positives**—Model predicts positive, actually positive. This is good.
- **Precision**—What proportion of positive predictions were actually correct? A model that produces no false positives has a precision of 1.0.
- **Recall**—What proportion of actual positives were predicted correctly ? A model that produces no false negatives has a recall of 1.0.
- **F1 score**—A combination of precision and recall. The closer to 1.0, the better.
- **Receiver operating characteristic (ROC) curve & Area under the curve (AUC)**—The ROC curve is a plot comparing true positive and false positive rate. The AUC metric is the area under the ROC curve. A model whose predictions are 100% wrong has an AUC of 0.0, one whose predictions are 100% right has an AUC of 1.0.

For regression problems (where you want to predict a number), you'll want to minimise the difference between what your model predicts and what the actual value is. If you're trying to predict the price a house will sell for, you'll want your model to get as close as possible to